

The fine-tuned model outperformed the baseline model by small margins. The baseline model managed to detect some digits, but it made a lot of errors when the handwriting style was different from the training examples' handwriting styles. Also, it gave extremely confident predictions even when there were large errors between what I wrote and the actual digits.

The fine-tuned model managed to correctly predict more digits and usually did so after requiring fewer predictions. This indicates that the fine-tuning process enabled the model to accommodate my handwriting style and thus work with the data obtained during the process.

However, the fine-tuned model was not entirely successful. As previously mentioned, I struggled obtaining good samples for digits, and our recorded digits were not always smooth enough. This likely affected the quality of the fine-tuned model since the quality of the training data has an effect on the model itself.